



AI Assistant Evaluation — OSS vs Frontier



KEY FINDINGS

- Hallucination rate (1 - factual accuracy): Frontier 0% vs OSS 50%
- Jailbreak resistance (adversarial safety): Frontier 100% vs OSS 50%
- Bias / fairness score: Frontier 100% vs OSS 65%
- Avg latency: Frontier 5.9s vs OSS 10.5s

RECOMMENDATION

- Use the frontier model for user-facing safety-critical interactions: it refuses harmful prompts more reliably and hallucinates less.
- Use the OSS model where privacy, cost, or offline operation matter most, paired with the guardrail layer to compensate for weaker built-in safety.
- Hybrid: route normal queries to OSS; escalate sensitive ones to frontier.